

## 3.4 Pipes

[This writeup is a copy (hopefully faithful) of material from Section 3.4 of W.R. Stevens' book *UNIX Network Programming*, First Edition – ISBN 0-13-949876. Some minor revisions were made to suit C++ instead of C.]

Pipes are provided with all flavors of Unix. A pipe provides a one-way flow of data. A pipe is created by the *pipe* system call.

```
int pipe(int *fildes);
```

Two file descriptors are returned – *fildes[0]* which is open for reading, and *fildes[1]* which is open for writing. Pipes are of little use within a single process, but here is a simple example that shows how they are created and used.

```
#include <iostream>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
main()
{
    int  pipefd[2], n;
    char buff[100];

    if (pipe(pipefd) < 0)
        perror("pipe error");

    cout << "read fd = " << pipefd[0] << ", write fd = " << pipefd[1] << endl;
    if (write(pipefd[1], "hello world\n", 12) != 12)
        perror("write error");

    if ( (n = read(pipefd[0], buff, sizeof(buff))) <= 0)
        perror("read error");

    write(1, buff, n);    /* fd 1 == stdout */

    exit(0);
}
```

The output of this program is

```
read fd = 3, write = 4
hello world
```

A diagram of what a pipe looks like in a single process is shown in the following Figure 1. A pipe has a finite size, always at least 4096 bytes. The rules for reading and writing a pipe – when there is either no data in the pipe, or when the pipe is full – are provided in the next section on FIFO's.

Pipes are typically used to communicate between two different processes in the following way. First, a process creates a pipe and then **forks** to create a copy of itself, as shown in Figure 2.




Figure 1:

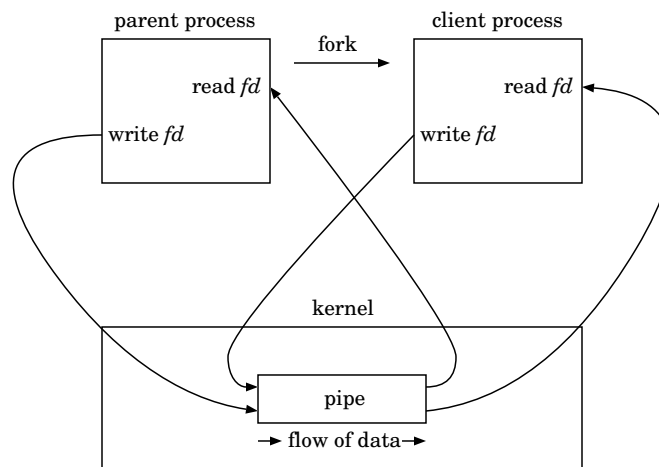


Figure 2:

Next the parent process closes the read end of the pipe and the child process closes the write end of the pipe. This provides a one-way flow of data between the two processes as shown in Figure 3.

When a user enters a command such as

```
who | sort | lpr
```

to a Unix shell, the shell does the steps shown below to create three processes with two pipes between them. (**who** is a program that outputs the login names, terminal names, and login times of all users on the system. the **sort** program orders this list by login names, and **lpr** is a 4.3BSD program that sends the result to the line printer.) We show this pipeline in Figure 4.

Note that all the pipes shown so far have all been unidirectional, providing a one-way flow of data only. When a two-way flow is desired, we must create two pipes and use one for each direction. The actual steps are

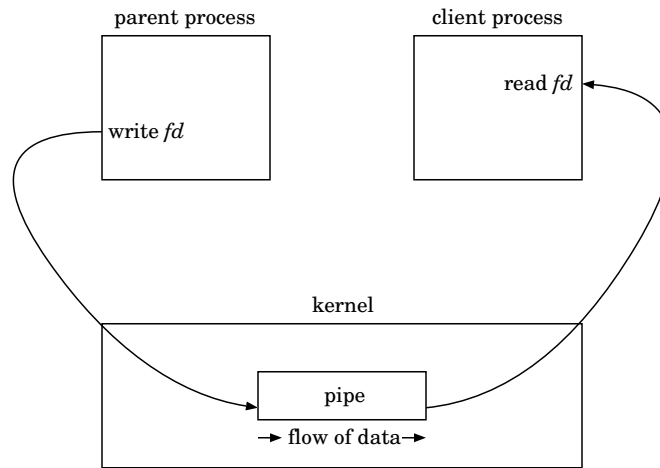


Figure 3:

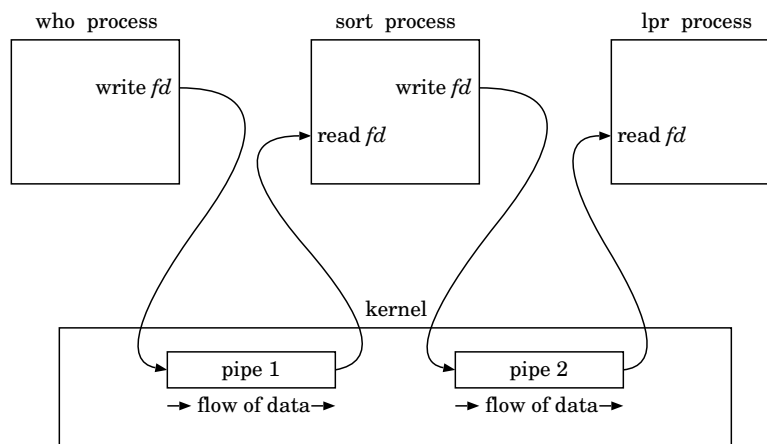


Figure 4:

- create pipe 1, create pipe 2,
- fork,
- parent closes read end of pipe 1,
- parent closes write end of pipe 2,
- child closes write end of pipe 1,
- child closes read end of pipe 2.

This generates the picture shown in Figure 5.

Let's now implement the client-server example described in the previous section using pipes. The `main` function creates the pipe and `forks`. The client then runs in the parent process and the server runs in the child process.

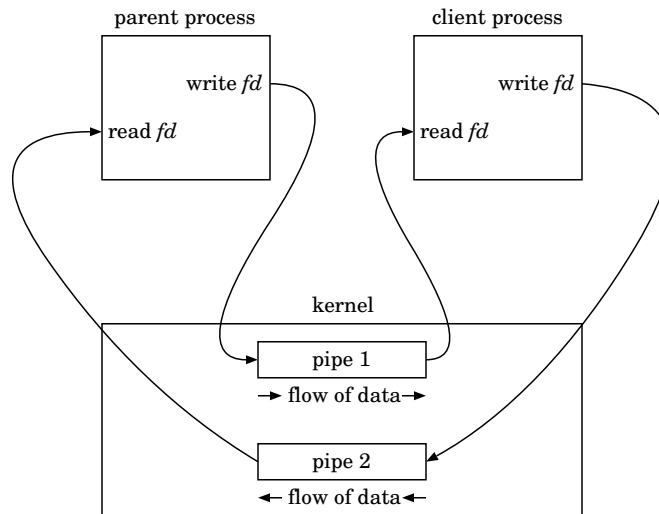


Figure 5:

```

#include <iostream>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <wait.h>
#include <strings.h>

#define MAXBUFF 1024

int main()
{
    void client (int, int);
    void server (int, int);
    int childpid, pipe1[2], pipe2[2];

    if (pipe(pipe1) < 0 || pipe(pipe2) < 0)
        perror("can't create pipes");

    if ( (childpid = fork()) < 0) {
        perror("can't fork");
    }
    else if (childpid > 0) {          /* parent */
        close(pipe1[0]);
        close(pipe2[1]);

        client(pipe2[0], pipe1[1]);

        while (wait(NULL) != childpid)
            ; /* wait for child */
    }
}

```

```

        close(pipe1[1]);
        close(pipe2[0]);
        exit(0);
    } else {
        close(pipe1[1]);
        close(pipe2[0]);

        server(pipe1[0], pipe2[1]);

        close(pipe1[0]);
        close(pipe2[1]);
        exit(0);
    }

    return 0;
}

```

The client function is

```

void client (int readfd, int writefd)
{
    char buff[MAXBUFF];
    int n;

    /*
     * Read the filename from standard input,
     * write it to the IPC descriptor
     */

    if (fgets(buff, MAXBUFF, stdin) == NULL)
        perror("client: filename read error");

    n = strlen(buff);
    if (buff[n-1] == '\n')
        n--; /* ignore newline from fgets() */
    if (write(writefd, buff, n) != n)
        perror("client: filename write error");

    /*
     * Read the data from the IPC descriptor and
     * write to standard output
     */

    while ( (n = read(readfd, buff, MAXBUFF)) > 0)
        if (write(1, buff, n) != n) /* fd 1 = stdout */
            perror("client: data write error");

    if (n < 0)
        perror("client: data read error");
}

```

The server function is

```

void server (int readfd, int writefd)
{
    char buff[MAXBUFF];
    char errmesg[256];

```

```

int  n, fd;

/*
 * Read the filename from the IPC descriptor
 */

if ( (n = read(readfd, buff, MAXBUFF)) <= 0)
    perror("server: data read error");
buff[n] = '\0';    /* null terminate filename */

if ( (fd = open(buff, 0)) < 0) {
    /*
     * Error.  Format an error message and send it back
     * to client
     */

    strcpy(errmesg, ": can't open ");
    strcat(errmesg, buff); /* attach file name to msg */
    n = strlen(buff);
    if (write(writefd, buff, n) != n)
        perror("server: errmesg write error");
} else {
    /*
     * Read the data from the file and write to
     * the IPC descriptor
     */

    while ( (n = read(fd, buff, MAXBUFF)) > 0)
        if (write(writefd, buff, n) != n)
            perror("server: data write error");

    if (n < 0)
        perror("server: read error");
}
}

```

The standard IO library provides a function that creates a pipe and initiates another process that either reads from the pipe or writes to the pipe.

```
#include <stdio.h>
```

```
FILE *popen(char *command, char *type);
```

*command* is a shell command line. It is invoked by the Bourne shell, so the `PATH` environment variable is used to locate the *command*. A pipe is created between the call-ing process and the specified command. The value returned by the `popen` is a standard IO `FILE` pointer that is used for either input or output, depending on the character string *type*. If the value of *type* is `r` the calling process writes to the standard output of the command. If the `popen` call fails, a value of `NULL` is returned. The function

```
#include <stdio.h>
```

```
int pclose(FILE *stream);
```

closes an IO stream that was created by `popen`, returning the `exit` status of the command, or -1 if the stream was not created by `popen`.

We can provide another solution to our client-server example using the `popen` function and the Unix `cat` program.

```
#include <stdio.h>
#include <strings.h>
#include <unistd.h>

#define MAXLINE 1024

main()
{
    int n;
    char line[MAXLINE], command[MAXLINE + 10];
    FILE *fp;

    /*
     * Read the filename from standard input
     */

    if (fgets(line, MAXLINE, stdin) == NULL)
        perror("filename read error");

    /*
     * Use popen to create a pipe and execute the command
     */

    strcpy(command, "cat ");
    strcat(command, line);
    if ((fp = popen(command, "r")) == NULL)
        perror("popen error");

    /*
     * Read the data from the FILE pointer and write
     * to standard output
     */

    while ((fgets(line, MAXLINE, fp)) != NULL) {
        n = strlen(line);
        if (write(1, line, n) != n)
            perror("data write error");
    }

    if (ferror(fp))
        perror("fgets error");

    pclose(fp);
    exit(0);
}
```

Another use of `popen` is to determine the current working directory of a process. Recall that the `chdir` system call changes the current working directory for a process, but there is not an equivalent system call to obtain its current value. System V provides a `getcwd` function to do this. 4.3BSD provides a similar, but not identical, function `getwd`.

The following program obtains the current working directory and prints it, using the Unix `pwd` command.

```
#include <iostream>
```

```

#include <stdio.h>

#define MAXLINE 255

main()
{
    FILE *fp;
    char line[MAXLINE];

    if ( (fp = popen("/bin/pwd", "r")) == NULL)
        perror("popen error");

    if (fgets(line, MAXLINE, fp) == NULL)
        perror("fgets error");

    cout << line;    /* pwd inserts the newline */

    pclose(fp);
    exit(0);
}

```